

Tuning Machine Learning Algorithms with mlr3

mlr3tuning

Department of Statistics – LMU Munich



Intro

TUNING

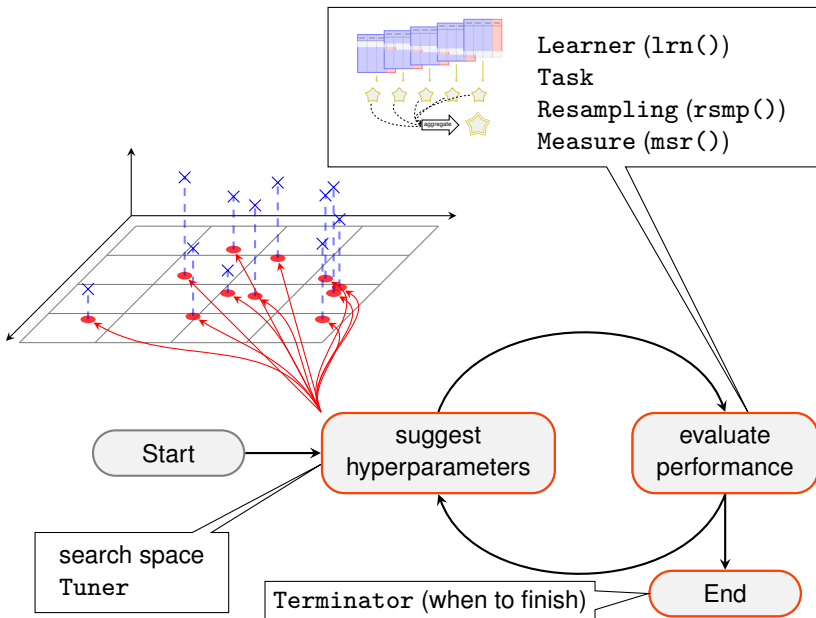
- Behavior of most methods depends on *hyperparameters*
 - We want to choose them so our algorithm performs well
 - Good hyperparameters are data-dependent
- ⇒ We do *black box optimization* (“Try stuff and see what works”)

Tuning toolbox for mlr3:

```
library("bbotk")  
library("mlr3tuning")
```


Tuning

TUNING




Tuning in mlr3

OBJECTS IN TUNING



SEARCH SPACE



```
ParamSet$new(list(param1, param2, ...))
```

Numerical parameter ParamDbl\$new(id, lower, upper)

Integer parameter ParamInt\$new(id, lower, upper)

Discrete parameter ParamFct\$new(id, levels)

Logical parameter ParamLgl\$new(id)

Untyped parameter ParamUty\$new(id)

```
library("paradox")
searchspace_knn = ParamSet$new(list(
  ParamInt$new("k", 1, 20)
))
```


TERMINATION

- Tuning needs a *termination condition*: when to finish
- Terminator class
- `mlr_terminators` dictionary, `trm()` short form

- `as.data.table(mlr_terminators)`

```
#> Key: <key>
#>           key
#>      <char>
#> 1:      clock_time
#> 2:           combo
#> 3:           evals
#> 4:           none
#> 5:    perf_reached
#> 6:         run_time
#> 7:        stagnation
#> 8: stagnation_batch
```

- `trm("evals", n_evals = 20)`

```
#> <TerminatorEvals>
#> * Parameters: n_evals=20
```

TUNING METHOD

- need to choose a *tuning method*
- Tuner class
- `mlr_tuners` dictionary, `tnr()` short form

- `as.data.table(mlr_tuners)`

```
#> Key: <key>
#>           key
#>      <char>
#> 1: design_points
#> 2:      gensa
#> 3:  grid_search
#> 4:      nloptr
#> 5: random_search
```

TUNING METHOD


- load Tuner with `tnr()`, set parameters

- `gsearch = tnr("grid_search", resolution = 3)`

```
print(gsearch)
#> <TunerGridSearch>
#> * Parameters: resolution=3, batch_size=1
#> * Parameter classes: ParamLgl, ParamInt, ParamDbl, ParamFct
#> * Properties: dependencies, single-crit, multi-crit
#> * Packages: -
```

- common parameter `batch_size` for parallelization

CALLING THE TUNER



```
inst = TuningInstanceSingleCrit$new(  
  tsk("iris"), lrn("classif.kknn", kernel="rectangular"),  
  rsmp("holdout"), msr("classif.ce"),  
  searchspace_knn, trm("none")  
)
```

CALLING THE TUNER

```
gsearch$optimize(inst)
```

```
#> INFO [10:23:07.146] Starting to optimize 1 parameter(s) with '<Optim
#> INFO [10:23:07.168] Evaluating 1 configuration(s)
#> INFO [10:23:07.687] Result of batch 1:
#> INFO [10:23:07.688]      k classif.ce      resample_result
#> INFO [10:23:07.688] 10      0.04 <ResampleResult[18]>
#> INFO [10:23:07.689] Evaluating 1 configuration(s)
#> INFO [10:23:07.783] Result of batch 2:
#> INFO [10:23:07.785]      k classif.ce      resample_result
#> INFO [10:23:07.785] 1      0.06 <ResampleResult[18]>
#> INFO [10:23:07.786] Evaluating 1 configuration(s)
#> INFO [10:23:07.814] Result of batch 3:
#> INFO [10:23:07.815]      k classif.ce      resample_result
#> INFO [10:23:07.815] 20      0.08 <ResampleResult[18]>
#> INFO [10:23:07.818] Finished optimizing after 3 evaluation(s)
#> INFO [10:23:07.819] Result:
#> INFO [10:23:07.819]      k learner_param_vals x_domain classif.ce
#> INFO [10:23:07.819] <num>      <list>      <list>      <num>
#> INFO [10:23:07.819]      10      <list[2]> <list[1]>      0.04
#>      k learner_param_vals x_domain classif.ce
#>      <num>      <list>      <list>      <num>
#> 1:      10      <list[2]> <list[1]>
```

TUNING RESULTS


```
gsearch = tnr("grid_search", resolution = 20)

inst = TuningInstanceSingleCrit$new(
  tsk("iris"), lrn("classif.kknn", kernel="rectangular"), rsmp("holdout"),
  msr("classif.ce"), searchspace_knn, trm("none"))

gsearch$optimize(inst)

#>      k learner_param_vals  x_domain classif.ce
#>   <num>         <list>    <list>      <num>
#> 1:    14         <list[2]> <list[1]>      0.02

ggplot(inst$archive$data(unnest = "x_domain"),
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```



RECAP

```
inst = TuningInstanceSingleCrit$new(  
  tsk("iris"), lrn("classif.kknn", kernel="rectangular"),  
  rsmp("holdout"), msr("classif.ce"),  
  searchspace_knn, trm("evals", n_evals = 2)  
)  
  
gsearch = tnr("grid_search", resolution = 3)  
  
gsearch$optimize(inst)  
  
#>      k learner_param_vals  x_domain classif.ce  
#>    <num>          <list>    <list>      <num>  
#> 1:      1          <list[2]> <list[1]>      0.06
```

Parameter Transformation

PARAMETER TRANSFORMATION

- Sometimes we do not want to sample evenly from a range
- $k = 1$ vs. $k = 2$ probably more interesting than $k = 101$ vs. $k = 102$

⇒ Transformations

- Part of ParamSet

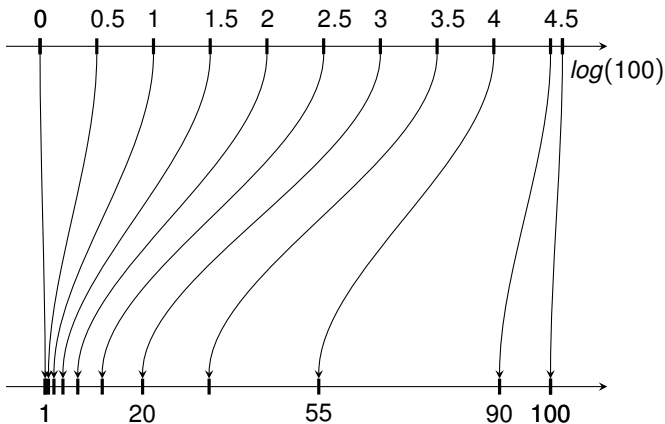
Example:

- 1 sample from $\log(1) \dots \log(100)$ (`k_before_trafo`)
- 2 transform by $\exp()$ in `trafo` function
- 3 don't forget to round (k must be integer)

```
searchspace_knn_trafo = ParamSet$new(list(  
  ParamDbl$new("k_before_trafo", log(1), log(50))  
))  
searchspace_knn_trafo$trafo = function(x, param_set) {  
  return(list(k = round(exp(x$k_before_trafo))))  
}
```

PARAMETER TRANSFORMATION

What is our transformation doing?




PARAMETER TRANSFORMATION

Tuning again...

```
inst$result  
  
#>      k_before_trafo learner_param_vals  x_domain classif.ce  
#>              <num>              <list>    <list>    <num>  
#> 1:                0          <list[2]> <list[1]>    0.02
```

PARAMETER TRANSFORMATION

```
ggplot(inst$archive$data(unnest = "x_domain"),  
  aes(x = x_domain_k, y = classif.ce)) + geom_line() + geom_point()
```



Nested Resampling

NESTED RESAMPLING

- Need to perform nested resampling to estimate tuned learner performance

⇒ Treat tuning as if it were a Learner!

- Training:

- 1 Tune model using (inner) resampling
- 2 Train final model with best parameters on all (i.e. outer resampling) data

- Predicting: Just use final model

- **AutoTuner**

```
optlrn = AutoTuner$new(lrn("classif.kknn", kernel="rectangular"),  
  rsmp("holdout"), msr("classif.ce"), searchspace_knn,  
  trm("none"), tnr("grid_search", resolution = 2))
```

NESTED RESAMPLING

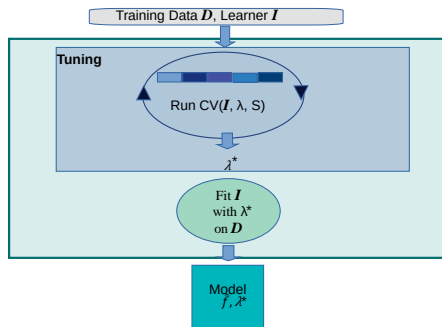

```
optlrn$train(tsk("iris"))
```

```
optlrn$model$learner
```

```
#> <LearnerClassifKNN:classif.kknn>
#> * Model: list
#> * Parameters: kernel=rectangular, k=1
#> * Packages: kknn
#> * Predict Type: response
#> * Feature types: logical, integer, numeric, factor, ordered
#> * Properties: multiclass, twoclass
```

```
ggplot(optlrn$model$tuning_instance$archive$data(),
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```

NESTED RESAMPLING



NESTED RESAMPLING

```
resample(tsk("iris"), optlrn, rsmp("holdout"))
```

```
#> <ResampleResult> of 1 iterations
```

```
#> * Task: iris
```

```
#> * Learner: classif.kknn.tuned
```


```
#> * Warnings: 0 in 0 iterations
```

```
#> * Errors: 0 in 0 iterations
```

NESTED RESAMPLING

```
result = resample(tsk("iris"), optlrn, rsmp("holdout"),  
  store_models = TRUE)
```

```
ggplot(result$learners[[1]]$  
  model$tuning_instance$archive$data(),  
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```



NESTED RESAMPLING

Aggregate performances of outer folds

```
result$aggregate()  
  
#> classif.ce  
#>      0.06
```

Retrieve inner tuning results

```
result$data$learner[[1]]$tuning_result  
  
#>      k learner_param_vals  x_domain classif.ce  
#>    <num>          <list>    <list>    <num>  
#> 1:      1          <list[2]> <list[1]>        0
```

Outro

TUNING WITH MLR3TUNING

Tuning a Learner

- 1 Construct a `TuningInstanceSingleCrit`
 - Task—the Data to tune over
 - Learner—the algorithm to tune
 - Resampling—the resampling method to use
 - Measure—how to evaluate performance
 - ParamSet—the search space, possibly with `trafo`
 - Terminator—when to quit
- 2 Create a Tuner
 - Usually using `tnr()`
 - May have some parameters, e.g. `batch_size`
- 3 Call `tuner$optimize()`

Nested Resampling

- 1 Construct an `AutoTuner`
 - Constructor takes all arguments of a `TuningInstanceSingleCrit` *except Task*
 - Also takes the Tuner as an argument
- 2 Use like a normal Learner in `resample()` and `benchmark()`